

Operating Systems
CS 3502 Sec 01 Spring 2025
Project 1

Matthew Pfleger
Kennesaw State University
mpfleger@students.kennesaw.edu
GitHub

February 28, 2025

Contents

1	Introduction	2
2	Environment Setup and Tool Usage	2
2.1	Development Environment	2
2.2	Tools and Programming Language	2
2.3	Resolving Issues	2
3	Implementation Details	2
3.1	Project A: Multi-Threading	2
3.1.1	Phase 1	3
3.1.2	Phase 2	4
3.1.3	Phase 3	6
3.1.4	Phase 4	8
3.1.5	Final	11
3.2	Project B: Inter-Process Communication	13
3.2.1	Producer Function	14
3.2.2	Consumer Function	15
3.2.3	Pipe Flow	15
3.2.4	Final	16
4	Challenges and Solutions	16
4.1	Challenges	16
4.2	Solutions	16
5	Results and Outcomes	16
6	Reflection and Learning	16
7	References	17

1 Introduction

I have broken down this project into two sub-projects, Project A and Project B. Project A focuses on multi-threaded programming, and Project B focuses on inter-process communication (IPC).

In Project A I have created a multi-threaded Bank Account system, that demonstrates the essential concepts of concurrent programming. Each phase of Project A builds upon the last exhibiting the core concepts of thread creation, resource protection, mutex locking, and deadlock.

In Project B I have created a program to exhibit inter-process communication (IPC) in programming using pipes. This program involves two processes a producer, who writes data, and a consumer, who reads the data from the producer via a pipe.

2 Environment Setup and Tool Usage

One of the most important aspects of this project was being able to get hands-on experience working with various different tooling and environments for the first time. Here I will document all of the various mechanisms I employed to make this project sing.

2.1 Development Environment

For this project I decided to use Oracle's Virtual Box for my virtual machine implementation. The reason I chose Virtual Box over the other virtual machine tools is because of its user friendly nature.

For the Linux distribution I chose Ubuntu based off the recommendation of it being very beginner-friendly, and well-documented. This made the process of researching how to use these tools much easier for this project.

2.2 Tools and Programming Language

This project is coded in C++ language. One of the earliest challenges I faced in this project was deciding which language to use. My background experience is mostly in Python and Java, both of which were excluded from this project. So, I had to learn the C++ syntax from the beginning.

For the IDE I chose Visual Studio Code because of my familiarity with this application. One of the challenges faced during the installation was learning how to install apps on a Linux based system. I have never used Linux before this project, so even something as seemingly simple as that posed a challenge for my project completion.

Another important tool for this assignment was the use of this PDF writing tool \LaTeX to create this document. This is another tool I have had no prior experience with, so learning it was challenging but very rewarding.

2.3 Resolving Issues

Many of the issues I faced while setting up my development environment were easily solvable. Because the tools I used are so well documented there is a plethora of online guides that helped walk me through each process. I have referenced these guides in the reference section of this document.

3 Implementation Details

3.1 Project A: Multi-Threading

For this multi-threading project the goal is to build a program piece by over four distinct phases. Each phase builds upon the last and demonstrates an aspect of multithreaded programming. These aspects include threading, resource protection, mutex locks, and deadlock.

3.1.1 Phase 1

In this phase, the goal is to demonstrate basic thread operations. These operations include how to create threads that perform concurrent operations, demonstrate proper thread creation and management, and show basic thread safety principles.

```
1 //Bank Account Class
2 #include <iostream>
3 class BankAccount{
4     private:
5         double balance; //Shared resource between threads
6     public:
7         //BankAccount constructor
8         BankAccount(double initial_balance){
9             balance = initial_balance;
10        }
11        //BankAccount destructor
12        ~BankAccount(){}
13
14        //deposit method
15        bool deposit(double amount){
16            if(amount <= 0){
17                std::cout << "ERROR: deposit amount must be greater than 0\n";
18                return false;
19            }
20            else{
21                balance += amount;
22            }
23            return true;
24        }
25        //withdraw method
26        bool withdraw(double amount){
27            if(amount > balance){
28                std::cout << "ERROR: Insufficient Funds\n";
29                return false;
30            }
31            else{
32                balance -= amount;
33            }
34            return true;
35        }
36        //get_balance method
37        double get_balance(){
38            return balance;
39        }
40 };
```

As you can see from my code above, I have created a simple bank account class that contains a balance variable and includes methods to withdraw and deposit money and to return the current balance.

To demonstrate threads running concurrent operations I created two tasks. Task one will deposit money, and task two will withdraw money from an instance of the bank account class. To run this task concurrently I have created two threads to complete the tasks.

```
1 #include "BankAccount1.cpp"
2 #include <thread>
3
4 //Task one for depositing into bank account
5 void task1(BankAccount &account, double amount){
6     std::cout << "You are withdrawing $" << amount << "\n";
7     account.withdraw(amount);
8 }
9
10 //Task two for withdrawing from bank account
11 void task2(BankAccount &account, double amount){
```

```

12     std::cout << "You are depositing $" << amount << "\n";
13     account.deposit(amount);
14 }
15
16 int main(){
17     //initialize bank account with starting balance of $100
18     BankAccount pflegerAccount(100);
19
20     //Output balance before thread operations
21     std::cout << "Balance before thread operations: $" << pflegerAccount.
get_balance() << "\n";
22
23     std::thread thread1(task1, std::ref(pflegerAccount), 30); //Creation of
thread1
24     std::thread thread2(task2, std::ref(pflegerAccount), 22); //Creation of
thread1
25
26     //join both threads to ensure they finish before program terminates
27     thread1.join();
28     thread2.join();
29
30     //Output balance after thread operations
31     std::cout << "Balance after thread operations: $" << pflegerAccount.
get_balance() << "\n";
32
33     return 0; //return 0;
34 }

```

In the main method I initialized a personal bank account with a starting balance of 100.00 dollars. After creating the threads and assigning them a task to run, I made sure to use the '.join()' keyword on both to ensure that the code waits to terminate until both threads have completed their task.

The above code will output the following:

```

Balance before thread operations: $100
You are depositing $22
You are withdrawing $30
Balance after thread operations: $92

```

3.1.2 Phase 2

The primary focus for phase two is demonstrating resource protection for multi-threaded coding. In the previous phase, I built a code that had two threads running concurrent operations on a bank account object. However, because the threads are both modifying the bank account object concurrently, the issue of race conditions can occur, when multiple threads are trying to access the same resource simultaneously.

To demonstrate race conditions I created eight threads that will deposit one dollar into a bank account 100,000 times (an exaggerated case required to demonstrate race conditions).

```

1 #include "BankAccount1.cpp"
2 #include <thread>
3 #include <vector>
4
5 //Method to simulate multiple deposits
6 void multipleDeposits(BankAccount &account, double amount, int run_time){
7     for(int i = 0; i < run_time; i++){
8         account.deposit(amount);
9     }
10 }
11
12 int main(){

```

```

13 //Initialize bank account with a starting balance of $0
14 BankAccount pflegerAccount(0);
15
16 //Use vector to hold 8 threads to run multipleDeposit task
17 std::vector<std::thread> threads;
18 for(int i = 0; i < 8; i++){
19     threads.push_back(std::thread(multipleDeposits, std::ref(pflegerAccount),
20     1, 100000));
21 }
22
23 //join all of the threads
24 for(auto&t : threads){
25     t.join();
26 }
27
28 //Output the results
29 std::cout << "Expected Balance: $800000\n";
30 std::cout << "Actual Balance: $" << pflegerAccount.get_balance() << "\n";
31
32 return 0;
33 }

```

For this code we would expect after execution for the balance to equal 800,000. However, due to race conditions our output will behave unpredictably. After running the code three times here are the outputs.

```

Output #1:
Expected Balance: $800000
Actual Balance: $796338

```

```

Output #2:
Expected Balance: $800000
Actual Balance: $709662

```

```

Output #3:
Expected Balance: $800000
Actual Balance: $718316

```

To solve this problem the BankAccount class from above builds on my code from phase one, but this time adding the necessary locking mechanisms required for resource protection. In the methods defined in the bank account class I included the addition of the mutex locks. The thread accesses the lock in order to access the critical section of the code.

```

1 //Bank Account Class
2 #include <iostream>
3 #include <mutex>
4 class BankAccount{
5     private:
6         double balance; //Shared resource between threads
7         std::mutex mtx;
8     public:
9         //BankAccount constructor
10        BankAccount(double initital_balance){
11            balance = initital_balance;
12        }
13        //BankAccount destructor
14        ~BankAccount(){}
15
16        //deposit method
17        bool deposit(double amount){
18            std::unique_lock<std::mutex> lock(mtx);

```

```

19         if(amount <= 0){
20             std::cout << "ERROR: deposit amount must be greater than 0\n";
21             return false;
22         }
23         else{
24             balance += amount;
25         }
26         return true;
27     }
28     //withdraw method
29     bool withdraw(double amount){
30         std::unique_lock<std::mutex> lock(mtx);
31         if(amount > balance){
32             std::cout << "ERROR: Insufficient Funds\n";
33             return false;
34         }
35         else{
36             balance -= amount;
37         }
38         return true;
39     }
40     //get_balance method
41     double get_balance(){
42         std::unique_lock<std::mutex> lock(mtx);
43         return balance;
44     }
45 };

```

After the additions of mutex locks, the output we get on our test code behaves as we expect.

Output #1:
 Expected Balance: \$800000
 Actual Balance: \$800000

Output #2:
 Expected Balance: \$800000
 Actual Balance: \$800000

Output #3:
 Expected Balance: \$800000
 Actual Balance: \$800000

3.1.3 Phase 3

The objective of phase three is to demonstrate deadlock creation. Deadlock occurs when two or more threads are trying to access a resource being locked by the other. When this happens the program is stuck endlessly waiting.

```

1 //Bank Account Class
2 #include <iostream>
3 #include <mutex>
4 #include <chrono>
5 #include <thread>
6 class BankAccount{
7     private:
8         double balance; //Shared resource between threads
9         std::mutex mtx;
10    public:
11        //BankAccount constructor
12        BankAccount(double initital_balance){
13            balance = initital_balance;

```

```

14     }
15     //BankAccount destructor
16     ~BankAccount(){}
17
18     //deposit method
19     bool deposit(double amount){
20         std::unique_lock<std::mutex> lock(mtx);
21         if(amount <= 0){
22             std::cout << "ERROR: deposit amount must be greater than 0\n";
23             return false;
24         }
25         else{
26             balance += amount;
27         }
28         return true;
29     }
30     //withdraw method
31     bool withdraw(double amount){
32         std::unique_lock<std::mutex> lock(mtx);
33         if(amount > balance){
34             std::cout << "ERROR: Insufficient Funds\n";
35             return false;
36         }
37         else{
38             balance -= amount;
39         }
40         return true;
41     }
42     //get_balance method
43     double get_balance(){
44         std::unique_lock<std::mutex> lock(mtx);
45         return balance;
46     }
47     static void transfer (BankAccount& sendingAccount, BankAccount&
receivingAccount, double amount){
48
49
50         //Check if sending account and recieving account are the same
51         if(&sendingAccount == &receivingAccount){
52             std::cout << "ERROR: Can not transfer between the same account\n";
53             return;
54         }
55
56         BankAccount *sender = &sendingAccount; //pointer to sender account
57         BankAccount *receiver = &receivingAccount; //pointer to receiver
account
58
59         std::unique_lock<std::mutex> lock1(sender->mtx);
60
61         std::this_thread::sleep_for(std::chrono::milliseconds(100000)); //
simulate a delay to demonstarte deadlock
62
63         std::unique_lock<std::mutex> lock2(receiver->mtx);
64
65         if(sender->balance >= amount){
66             sender->balance -= amount;
67             receiver->balance += amount;
68         }
69         else{
70             std::cout << "ERROR: Insufficient funds for transfer\n";
71         }
72     }
73 }
74 };

```

In the code above when we try to test our new 'transfer' method in the BankAccount class we

will experience deadlock. I included a delay in the method to demonstrate a potential deadlock scenario.

```
1 #include "BankAccount3.cpp"
2 #include <thread>
3
4 //task1 to transfer from pflegerAccount to johnsonAccount
5 void task1(BankAccount& acc1, BankAccount& acc2) {
6     BankAccount::transfer(acc1, acc2, 100);
7 }
8 //task2 to transfer from johnsonAccount to pflegerAccount
9 void task2(BankAccount& acc1, BankAccount& acc2) {
10     BankAccount::transfer(acc2, acc1, 50);
11 }
12
13 int main(){
14
15     //Creation of two test bank accounts
16     BankAccount pflegerAccount(500);
17     BankAccount johnsonAccount(500);
18
19     //Thread creation
20     std::thread thread1(task1, std::ref(pflegerAccount), std::ref(johnsonAccount))
21     ;
22     std::thread thread2(task2, std::ref(pflegerAccount), std::ref(johnsonAccount))
23     ;
24
25     //Thread join
26     thread1.join();
27     thread2.join();
28
29     //Outout the results
30     std::cout << "Pfleger Account Balance: $" << pflegerAccount.get_balance() << "
31     \n";
32     std::cout << "Johnson Account Balance: $" << johnsonAccount.get_balance() << "
33     \n";
34 }
```

When the tasks are run the first account will lock, then after the delay the second account will lock then each task will endlessly wait for the other to unlock, never being able to finish the required task.

3.1.4 Phase 4

In this phase I will demonstrate multiple ways to solve the deadlock issues that we created in phase three.

```
1 //Bank Account Class
2 #include <iostream>
3 #include <mutex>
4 #include <chrono>
5 class BankAccount{
6     private:
7         double balance; //Shared resource between threads
8         int id; //Unqiue id for each account
9         std::mutex mtx; //mutex lock for security
10    public:
11        //BankAccount constructor
12        BankAccount(double initital_balance, int user_id){
13            balance = initital_balance;
14            id = user_id;
15        }
16        //BankAccount destructor
17        ~BankAccount(){}
18 }
```



```

19 //deposit method
20 bool deposit(double amount){
21     //Attempt to lock mutex (if fails mutex already locked) avoid deadlock
22     if(mtx.try_lock()){
23         if(amount <= 0){
24             std::cout << "ERROR: deposit amount must be greater than 0\n";
25             return false;
26         }
27         else{
28             balance += amount;
29             mtx.unlock();//Explicitly unlock
30         }
31     }
32     return true;
33 }
34 //withdraw method
35 bool withdraw(double amount){
36     //Attempt to lock mutex (if fails mutex already locked) avoid deadlock
37     if(mtx.try_lock()){
38         if(amount > balance){
39             std::cout << "ERROR: Insufficient Funds\n";
40             return false;
41         }
42         else{
43             balance -= amount;
44         }
45         mtx.unlock();//Explicitly unlock
46     }
47     return true;
48 }
49
50 //get_balance method
51 double get_balance(){
52     //Attempt to lock mutex (if fails mutex already locked) avoid deadlock
53     if(mtx.try_lock()){
54         return balance;
55         mtx.unlock();//Explicitly unlock
56     }
57     return 0;
58 }
59 int get_id(){
60     return id;
61 }
62 static void transfer(BankAccount& sendingAccount, BankAccount&
receivingAccount, double amount) {
63     // Check if sending and receiving accounts are the same
64     if (&sendingAccount == &receivingAccount) {
65         std::cout << "ERROR: Cannot transfer between the same account\n";
66         return;
67     }
68
69     // Determine the lock order without affecting sender/receiver roles
70     BankAccount* first = &sendingAccount;
71     BankAccount* second = &receivingAccount;
72
73     if (sendingAccount.get_id() > receivingAccount.get_id()) {
74         std::swap(first, second);
75     }
76
77     std::unique_lock<std::mutex> lock1(first->mtx, std::defer_lock);
78     std::unique_lock<std::mutex> lock2(second->mtx, std::defer_lock);
79
80     // Lock both mutexes simultaneously to avoid deadlocks
81     std::lock(lock1, lock2);
82

```

```

83         // Perform the transfer correctly using the original sender/receiver
        references
84         if (sendingAccount.balance >= amount) {
85             sendingAccount.balance -= amount;
86             receivingAccount.balance += amount;
87             std::cout << "Transfer complete\n";
88         } else {
89             std::cout << "ERROR: Insufficient funds for transfer\n";
90         }
91     }
92
93 };

```

In the updated code, I have added an id variable for each instance of the bank account to use in consistent lock ordering. Inside, the getbalance, deposit, and withdraw methods I have included a timeout mechanisms to ensure the prevention of deadlock. Inside of the transfer method I have included a mechanism for lock ordering, ensuring the code will always lock the lower id first. Both account locks are locked at the same time to prevent deadlock.

```

1  #include "BankAccount4.cpp"
2  #include <thread>
3
4  //task1 to transfer from pflegerAccount to johnsonAccount
5  void task1(BankAccount& acc1, BankAccount& acc2) {
6      BankAccount::transfer(acc1, acc2, 100);
7  }
8  //task2 to transfer from johnsonAccount to pflegerAccount
9  void task2(BankAccount& acc1, BankAccount& acc2) {
10     BankAccount::transfer(acc2, acc1, 50);
11 }
12
13 int main(){
14
15     //Creation of two test bank accounts
16     BankAccount pflegerAccount(500,2);
17     BankAccount johnsonAccount(500,1);
18
19     //Thread creation
20     std::thread thread1(task1, std::ref(pflegerAccount), std::ref(johnsonAccount))
21     ;
22     std::thread thread2(task2, std::ref(pflegerAccount), std::ref(johnsonAccount))
23     ;
24
25     //Thread join
26     thread1.join();
27     thread2.join();
28
29     //Outout the results
30     std::cout << "Pfleger Account Balance: $" << pflegerAccount.get_balance() << "
    \n";
    std::cout << "Johnson Account Balance: $" << johnsonAccount.get_balance() << "
    \n";
}

```

With the implementations added to the code, we can now complete the inter-account transfers concurrently! Without the threat of deadlocks our code outputs the following.

```

Transfer complete
Transfer complete
Pfleger Account Balance: $450
Johnson Account Balance: $550

```

3.1.5 Final

In this project I created a bank account class that has the ability to perform basic thread operations, displays resource protection, and effectively prevents the creation of deadlock. Displayed below is the final version of my code built piece by piece during each phase of development.

```
1 //Bank Account Class
2 #include <iostream>
3 #include <mutex>
4 #include <chrono>
5 #include <thread>
6 #include <vector>
7 class BankAccount{
8     private:
9         double balance; //Shared resource between threads
10        int id; //Unqiue id for each account
11        std::mutex mtx; //mutex lock for security
12    public:
13        //BankAccount constructor
14        BankAccount(double initital_balance, int user_id){
15            balance = initital_balance;
16            id = user_id;
17        }
18        //BankAccount destructor
19        ~BankAccount(){}
20
21        //deposit method
22        bool deposit(double amount){
23            //Attempt to lock mutex (if fails mutex already locked) avoid deadlock
24            if(mtx.try_lock()){
25                if(amount <= 0){
26                    std::cout << "ERROR: deposit amount must be greater than 0\n";
27                    return false;
28                }
29                else{
30                    balance += amount;
31                    mtx.unlock();//Explicitly unlock
32                }
33            }
34            return true;
35        }
36        //withdraw method
37        bool withdraw(double amount){
38            //Attempt to lock mutex (if fails mutex already locked) avoid deadlock
39            if(mtx.try_lock()){
40                if(amount > balance){
41                    std::cout << "ERROR: Insufficient Funds\n";
42                    return false;
43                }
44                else{
45                    balance -= amount;
46                }
47                mtx.unlock();//Explicitly unlock
48            }
49            return true;
50        }
51
52        //get_balance method
53        double get_balance(){
54            //Attempt to lock mutex (if fails mutex already locked) avoid deadlock
55            if(mtx.try_lock()){
56                return balance;
57                mtx.unlock();//Explicitly unlock
58            }
59            return 0;
```

```

60     }
61     int get_id(){
62         return id;
63     }
64     static void transfer(BankAccount& sendingAccount, BankAccount&
receivingAccount, double amount) {
65         // Check if sending and receiving accounts are the same
66         if (&sendingAccount == &receivingAccount) {
67             std::cout << "ERROR: Cannot transfer between the same account\n";
68             return;
69         }
70
71         // Determine the lock order without affecting sender/receiver roles
72         BankAccount* first = &sendingAccount;
73         BankAccount* second = &receivingAccount;
74
75         if (sendingAccount.get_id() > receivingAccount.get_id()) {
76             std::swap(first, second);
77         }
78
79         std::unique_lock<std::mutex> lock1(first->mtx, std::defer_lock);
80         std::unique_lock<std::mutex> lock2(second->mtx, std::defer_lock);
81
82         // Lock both mutexes simultaneously to avoid deadlocks
83         std::lock(lock1, lock2);
84
85         // Perform the transfer correctly using the original sender/receiver
references
86         if (sendingAccount.balance >= amount) {
87             sendingAccount.balance -= amount;
88             receivingAccount.balance += amount;
89             std::cout << "Transfer complete\n";
90         } else {
91             std::cout << "ERROR: Insufficient funds for transfer\n";
92         }
93     }
94 };
95
96
97 int main(){
98     BankAccount account1(1356.88, 1234);
99     BankAccount account2(1245.87, 4321);
100    BankAccount account3(2234.12, 5678);
101    BankAccount account4(987.66, 8765);
102    BankAccount account5(500.91, 9876);
103
104    std::vector<std::thread> threads;
105    threads.push_back(std::thread(&BankAccount::transfer, std::ref(account1), std
::ref(account2), 100));
106    threads.push_back(std::thread(&BankAccount::transfer, std::ref(account2), std
::ref(account3), 200));
107    threads.push_back(std::thread(&BankAccount::transfer, std::ref(account3), std
::ref(account4), 300));
108    threads.push_back(std::thread(&BankAccount::transfer, std::ref(account4), std
::ref(account5), 400));
109    threads.push_back(std::thread(&BankAccount::transfer, std::ref(account5), std
::ref(account1), 500));
110
111    threads.push_back(std::thread(&BankAccount::deposit, std::ref(account1), 100))
;
112    threads.push_back(std::thread(&BankAccount::deposit, std::ref(account2), 200))
;
113    threads.push_back(std::thread(&BankAccount::deposit, std::ref(account3), 300))
;

```

```

114     threads.push_back(std::thread(&BankAccount::deposit, std::ref(account4), 400))
115     ;
116     threads.push_back(std::thread(&BankAccount::deposit, std::ref(account5), 500))
117     ;
118     threads.push_back(std::thread(&BankAccount::withdraw, std::ref(account1), 100)
119     );
120     threads.push_back(std::thread(&BankAccount::withdraw, std::ref(account2), 200)
121     );
122     threads.push_back(std::thread(&BankAccount::withdraw, std::ref(account3), 300)
123     );
124     threads.push_back(std::thread(&BankAccount::withdraw, std::ref(account4), 400)
125     );
126     threads.push_back(std::thread(&BankAccount::withdraw, std::ref(account5), 500)
127     );
128
129     for(auto& thread : threads){
130         thread.join();
131     }
132
133     std::cout << "Account 1 balance: $" << account1.get_balance() << std::endl;
134     std::cout << "Account 2 balance: $" << account2.get_balance() << std::endl;
135     std::cout << "Account 3 balance: $" << account3.get_balance() << std::endl;
136     std::cout << "Account 4 balance: $" << account4.get_balance() << std::endl;
137     std::cout << "Account 5 balance: $" << account5.get_balance() << std::endl;
138
139     return 0;
140 }

```

Included is a main method to test the BankAccount class on 15 different thread operations. Five multi account transfers, deposits, and withdrawals to test the integrity of my code. Here is the output after completion.

```

Transfer complete
Transfer complete
Transfer complete
Transfer complete
Transfer complete
Account 1 balance: $1756.88
Account 2 balance: $1145.87
Account 3 balance: $2134.12
Account 4 balance: $887.66
Account 5 balance: $400.91

```

3.2 Project B: Inter-Process Communication

For the inter-process communication (IPC) assignment, I chose to demonstrate this communication via pipes. The producer and consumer functions will communicate data through the pipe to be read out in the main method.

```

1 //Demonstrates interprocess communication using pipes
2 //Parent process creates a pipe and forks a child process
3
4 #include <iostream>
5 #include <unistd.h>
6 #include <cstring>
7 #include <sys/wait.h>
8
9 //Producer process
10 void producer(int pipe_fd[2]){
11     //Close read end of the pipe

```

```

12     close(pipe_fd[0]);
13     std::cout << "Producer: Writing message to the pipe\n";
14     //Message pointer
15     const char* message = "It's-a me, Mario!";
16     //Write message to the pipe
17     write(pipe_fd[1], message, strlen(message) + 1);
18     //Close write end of the pipe
19     close(pipe_fd[1]);
20 }
21
22 //Consumer process
23 void consumer(int pipe_fd[2]){
24     //Close write end of the pipe
25     close(pipe_fd[1]);
26     std::cout << "Consumer: Reading message from the pipe\n";
27     //Buffer to store the message
28     char buffer[256];
29     //Read message from the pipe
30     read(pipe_fd[0], buffer, sizeof(buffer));
31     //Print the message
32     std::cout << "Consumer has recieved the message!: " << buffer << std::endl;
33     //Close read end of the pipe
34     close(pipe_fd[0]);
35 }
36
37 //Main function to show interprocess communication using pipes
38 int main(){
39     int pipe_fd[2];
40
41     //Create a pipe
42     if(pipe(pipe_fd) == -1){ //Test for pipe creation failure
43         std::cerr << "Pipe creation failed!\n";
44         return 1;
45     }
46
47     pid_t pid = fork(); //Fork a child process
48
49     if(pid == -1){ //Test for fork failure
50         std::cerr << "Fork failed!\n";
51         return 1;
52     }
53
54     if(pid == 0){ //Child process
55         consumer(pipe_fd);
56     }else{ //Parent process
57         producer(pipe_fd);
58         wait(NULL); //Wait for the child process to finish
59     }
60
61     return 0;
62 }

```

3.2.1 Producer Function

In the producer function it takes an array of two file descriptors representing the pipe as its argument. It starts by closing the 'read' end of the pipe, because a producer will only write to the pipe. The function defines a pointer of the message then write the message to the 'write' end of the pipe. Finally, It will close the 'write' end to ensure resource integrity.

```

//Producer process
void producer(int pipe_fd[2]){
    //Close read end of the pipe
    close(pipe_fd[0]);

```

```

    std::cout << "Producer: Writing message to the pipe\n";
    //Message pointer
    const char* message = "It's-a me, Mario!";
    //Write message to the pipe
    write(pipe_fd[1], message, strlen(message) + 1);
    //Close write end of the pipe
    close(pipe_fd[1]);
}

```

3.2.2 Consumer Function

In the consumer function, just as in the producer function, it takes an array of two file descriptors representing the pipe as its argument. It starts by closing the 'write' end of the pipe, because a consumer will only read from the pipe. The function defines a buffer to store the read data, and reads the data from the producer function and outputs the received message. Then, to ensure resource integrity, the function closes the 'read' end of the pipe.

```

//Consumer process
void consumer(int pipe_fd[2]){
    //Close write end of the pipe
    close(pipe_fd[1]);
    std::cout << "Consumer: Reading message from the pipe\n";
    //Buffer to store the message
    char buffer[256];
    //Read message from the pipe
    read(pipe_fd[0], buffer, sizeof(buffer));
    //Print the message
    std::cout << "Consumer has recieved the message!: " << buffer << std::endl;
    //Close read end of the pipe
    close(pipe_fd[0]);
}

```

3.2.3 Pipe Flow

The main method of the code is where you can see the pipe flow being demonstrated. First, the main method initializes the array that holds the file descriptors for the pipe. Then, a pipe is created and check for any errors, then the pipe is forked into a parent, and child process each with their own process id. After the fork process the main method checks for any errors. If the process id equals 0 this indicates the consumer, and will run the consumer function. Conversely, if the process id equals 1 then the producer function will run. Before the code terminates it calls a wait to ensure the child process has time to finish before termination.

```

//Main function to show interprocess communication using pipes
int main(){
    int pipe_fd[2];

    //Create a pipe
    if(pipe(pipe_fd) == -1){ //Test for pipe creation failure
        std::cerr << "Pipe creation failed!\n";
        return 1;
    }
}

```

```

pid_t pid = fork(); //Fork a child process

if(pid == -1){ //Test for fork failure
    std::cerr << "Fork failed!\n";
    return 1;
}

if(pid == 0){ //Child process
    consumer(pipe_fd);
}else{ //Parent process
    producer(pipe_fd);
    wait(NULL); //Wait for the child process to finish
}

return 0;
}

```

3.2.4 Final

The program will output the following message provided by the producer function.

```

Consumer: Reading message from the pipe
Producer: Writing message to the pipe
Consumer has recieved the message!: It's-a me, Mario!

```

4 Challenges and Solutions

4.1 Challenges

One of the biggest challenges faced during this project was learning all of the new tools required to finish this assignment.

4.2 Solutions

Like previously mentioned I was able to overcome my lack of knowledge in the required tooling with the wealth of online guides that I have referenced in the reference section of this document.

5 Results and Outcomes

In this project I created a multi-threaded bank account system that can perform tasks on ten threads or more, with proper resource protection and deadlock prevention mechanisms. I also created a program that demonstrates inter-process communication (IPC) via pipes, a producer function passes a message to a consumer function that relays the message.

One of the main areas of growth for this project is combining the concepts I learned in multi-threading and inter-process communication (IPC). Potential integration of this could be creating a pipeline where threaded processes communicate via pipes.

6 Reflection and Learning

Throughout this process I learned a lot about multi-threading and inter-process communication. With learning how useful multi-threaded programming can be when writing complex programs that

have many processes running at once. Also, learning that it is important to be familiar with the resource protection functionalities, and how to prevent deadlock. It was fascinating to learn and develop a program that could communicate data via pipes.

But, not only that I learned about new tools that can help me in my future projects. Tools such as virtual machines, the Linux operating system, L^AT_EX PDF editor, and the C++ programming language.

7 References

References

- [1] CppNuts, “Threading in c++ — complete course,” 03 2021.
- [2] S. Mahapatra, “Multithreading in c++,” 01 2018.
- [3] GeeksForGeeks, “C tutorial - learn c programming language.”
- [4] freeCodeCamp.org, “Latex – full tutorial for beginners,” 03 2023.
- [5] GeeksForGeeks, “Inter process communication (ipc),” 01 2017.
- [6] T. Programmer, “How to install ubuntu 24.04 lts in virtualbox 2024,” 08 2024.
- [7] freeCodeCamp.org, “Git and github for beginners - crash course,” 05 2020.